

Chapter 5: Distributions

DSCC 462

Computational Introduction to Statistics

Anson Kahng

Fall 2022

Random Variables

- **Random variable:** A variable that can take a number of different values and whose outcome is determined by chance
- **Discrete random variable:** A random variable whose possible outcomes are a list of discrete values (finite or countably infinite sample space)
 - Example: Coin flip (heads/tails)
- **Continuous random variable:** A random variable whose possible outcomes are any value in an interval (uncountable sample space)
 - Examples: Time required to run a mile

Notation

- Random variable: Uppercase letters (e.g., X , Y)
- Outcome of a random variable: Lowercase letters (e.g., x , y)
- Example: Let X = the number of surgeries a person has had
 - $\Pr(X = 1)$: Probability of having 1 surgery
 - $\Pr(X = x)$: Probability of having x surgeries

Probability Distribution

- **Probability distribution:** List of all possible values that a random variable can take, along with their corresponding probabilities
 - Discrete: Probability mass function (PMF)
 - Continuous: Probability density function (PDF)
- Let X be a random variable defined over sample space S_X
- For any $E \subseteq S_X$, we can define $p_X(E) = \Pr(X \in E)$

Discrete Probability Distribution

- For a discrete random variable X with sample space S_X , a probability mass function (PMF) p_X maps $x \in S_X$ to a number in $[0,1]$ such that:

$$0 \leq p_X(s) = \Pr(X = x) \leq 1$$

$$\sum_{x \in S_X} p_X(x) = \sum_{x \in S_X} \Pr(X = x) = 1$$

- The support S_X consists of all x for which $p_X(x) > 0$ (all achievable outcomes)

Discrete Probability Distribution: Example

- Setup: A fair coin is flipped 3 times. Let X be a random variable that counts the number of heads observed
- Fill in the following table:
- Probability distribution tables resemble relative frequency distribution tables: probability of each outcome is the relative frequency distribution of each outcome in a large number of trials

x	$\Pr(X = x)$
0	
1	
2	
3	

Discrete Probability Distribution: Example

- Setup: A fair coin is flipped 3 times. Let X be a random variable that counts the number of heads observed
- Fill in the following table:
- Probability distribution tables resemble relative frequency distribution tables: probability of each outcome is the relative frequency distribution of each outcome in a large number of trials

x	$\Pr(X = x)$
0	1/8
1	3/8
2	3/8
3	1/8

Continuous Probability Distribution

- Specify continuous probability distributions through a *density function*, $f(x)$
- Properties:

$$f(x) \geq 0, \text{ for all } x \in S_X \text{ (nonnegative density)}$$

$$\int f(x) dx = 1 \text{ (total probability is 1)}$$

- X is continuous iff there is a density function f_X such that the following holds:

$$\begin{aligned} \Pr(a \leq X \leq b) &= \int_a^b f_X(x) dx \\ &= \text{Area under } f \text{ between } a \text{ and } b \end{aligned}$$

- The support S_X consists of all x for which $f_X(x) > 0$

Normalization

- We must ensure that probability distributions sum / integrate to 1 (i.e., total probability must equal 1)
- **Normalization:** Scalar adjustment in order to ensure that $\Pr(S_X) = 1$
- If $g(x) > 0$ for all $x \in S_X$, then


$$\text{Discrete: } p(x) = \frac{g(x)}{\sum_{x^* \in S_X} g(x^*)}$$

$$\text{Continuous: } f(x) = \frac{g(x)}{\int_{x^* \in S_X} g(x^*) dx^*}$$

- *Normalization constant:* 1/denominator

Normalization: Example

- Suppose that we generally know that probability is distributed according to the following curve:



Normalization: Example

- We can generally define the shape of this curve as

$$g(x) = 1 - x^2, \quad -1 \leq x \leq 1$$

- Is this a proper density?
- What's the normalization constant?
- What is $f(x)$?

Normalization: Example

- We can generally define the shape of this curve as

$$g(x) = 1 - x^2, \quad -1 \leq x \leq 1$$

- Is this a proper density?

$$\int_{-1}^1 (1 - x^2) dx = \left(x - x^3/3 \right) \Big|_{-1}^1 = \frac{4}{3}$$

- What's the normalization constant?

$$f(x) = \frac{g(x)}{\int g(x) dx}$$

Multiply both sides by $\frac{3}{4}$ in order to make it integrate to 1

- What is $f(x)$?

$$f(x) = \frac{3}{4}(1 - x^2)$$

Cumulative Distribution Functions (CDFs)

- The **cumulative distribution function (CDF)** of random variable X is

$$F_X(x) = \Pr(X \leq x) \text{ for all } x \in (-\infty, \infty)$$

- If X is discrete with support S_X , then the CDF is defined as

$$F_X(x) = \Pr(X \leq x) = \sum_{u \in S_X: u \leq x} \Pr(X = u)$$

- If X is continuous, then the CDF is defined as

$$F_X(x) = \Pr(X \leq x) = \int_{-\infty}^x f_X(u) du$$

Cumulative Distribution Functions (CDFs)

- Consider the parabolic density



$$f(x) = \begin{cases} 0 & x \in (-\infty, -1) \\ \frac{3}{4}(1 - x^2) & x \in [-1, 1] \\ 0 & x \in (1, \infty) \end{cases}$$

- Over the range $x \in (-\infty, -1)$, we have $F(x) = 0$
- Over the range $x \in [-1, 1]$, we have

$$F(x) = \int_{-1}^x \frac{3}{4}(1 - u^2) du = -x^3/4 + 3x/4 + 1/2$$

f(x)

- Over the range $x \in (1, \infty)$, we have $F(x) = 1$



Quantiles and Percentiles

Percentiles and Quartiles

- Suppose that a student with an 85 on an exam scored higher than 72% of their classmates
- Then, $\Pr(X \leq 85) = 0.72$
- We say that $q = 85$ is the $p = 0.72$ quantile of this distribution (also called the 72nd percentile)

Quantiles and Percentiles

- More generally: For a random variable X , q is the p -quantile of X if

$$\Pr(X < q) \leq p \text{ and } \Pr(X > q) \leq 1 - p$$

- The quantile function of X is then defined as

$$Q(p) = \min\{x \in S_X : \Pr(X \leq x) \geq p\}$$


- If the CDF $F_X(x)$ is continuous and strictly increasing on S_X , then

$$Q(p) = F_x^{-1}(p)$$

- Although $Q(p)$ is uniquely defined, the p -quantile may not be unique

Quantiles and Percentiles: Example

- Consider the parabolic density, $f(x) = \frac{3}{4}(1 - x^2)$
parabola 抛物线
- What is the 0.25-quantile?



Quantiles and Percentiles: Example

- Consider the parabolic density, $f(x) = \frac{3}{4}(1 - x^2)$
- What is the 0.25-quantile?

$$Q(p) = \min\{x \in S_X : \Pr(X \leq x) \geq p\}$$

Want to find x such that $\Pr(X \leq x) = 0.25$

Solve for x :

$$\text{integrate}(f(x)) = 0.25$$


$$-x^3/4 + 3x/4 + 1/2 = 1/4$$

in r: `polyroot(c(-1,-3,0,1))` $x^3 - 3x - 1 = 0$

this function solves the following equation: $-1-3x+0x^2+x^3 = 0$ roots: $x = \{-1.53, -0.35, 1.879\}$

Because we know that $x \in [-1, 1]$, we have that the 0.25-quantile occurs at

$$x = -0.35$$



Summarizing Probability Distributions

- Many random variables can take a large number of values, so an explicit probability distribution may be quite complicated
- We can describe a probability distribution with measures of central tendency and dispersion
- *Population mean*: Average value that a random variable takes
- *Population variance*: Dispersion of the values relative to the population mean
- *Population standard deviation*: The square root of the population variance

Expected Value

The Expected value of X is the population mean of X

- **Expected value** of X , denoted $E(X)$, represents a theoretical average of an infinitely large sample
 - $E(X)$ is what we “expect” X to equal; the population mean of X
- We use the notation $\mu = \mu_X = E(X)$

Expected Value

- If X is a discrete random variable:

$$\mu_X = E(X) = \sum_{x \in S_X} x \cdot \Pr(X = x)$$

- If X is a continuous random variable:

$$\mu_X = E(X) = \int_{-\infty}^{\infty} x \cdot f_X(x) dx$$

- If c is a constant, then

$$E(c) = c$$

Linearity of Expectation

- For any random variables X and Y :

$$E(X + Y) = E(X) + E(Y)$$

- This holds even if X and Y are *not* independent
- In general, for random variables X_1, \dots, X_n :

$$E\left(\sum_{i=1}^n X_i\right) = \sum_{i=1}^n E(X_i)$$

Mean of a Random Variable: Example


- Q1: What is the expected number of heads when flipping a fair coin ($1/2$ H, $1/2$ T)?
- Q2: What is the expected number of heads when flipping an unfair coin ($2/3$ H, $1/3$ T)?
- Q3: What is the expected number of heads when flipping three fair coins ($1/2$ H, $1/2$ T) and two unfair coins ($2/3$ H, $1/3$ T)?

Mean of a Random Variable: Example

- Q1: What is the expected number of heads when flipping a fair coin (1/2 H, 1/2 T)?
Let X be the random variable representing the number of heads
 $E(X) = (1/2) \cdot 1 + (1/2) \cdot 0 = 1/2$
- Q2: What is the expected number of heads when flipping an unfair coin (2/3 H, 1/3 T)?
 $E(X) = (2/3) \cdot 1 + (1/3) \cdot 0 = 2/3$
- Q3: What is the expected number of heads when flipping three fair coins (1/2 H, 1/2 T) and two unfair coins (2/3 H, 1/3 T)?
 $E(X) = 3 \cdot (1/2) + 2 \cdot (2/3) = 3/2 + 4/3 = 17/6$ by linearity of expectation

Mean of a Random Variable: Example

- Consider the parabolic density, $f(x) = \frac{3}{4}(1 - x^2)$
- What is $E(X)$?




- Intuition: By symmetry, $E(X) = 0$

Mean of a Random Variable: Example

- Consider the parabolic density, $f(x) = \frac{3}{4}(1 - x^2)$
- What is $E(X)$?

$$\begin{aligned} E(X) &= \int_{-1}^1 (x) \frac{3}{4}(1 - x^2) dx = \int_{-1}^1 \frac{3x}{4} - \frac{3x^3}{4} dx \\ &= \frac{3x^2}{8} - \frac{3x^4}{16} \Big|_{-1}^1 \\ &= \frac{3}{8} - \frac{3}{16} - \left(\frac{3}{8} - \frac{3}{16} \right) \\ &= 0 \end{aligned}$$

Intuition: By symmetry, $E(X) = 0$



Variance

- The variance of X , denoted $var(X)$, measures the tendency of X to deviate from $E(X)$ and is defined as follows

$$\begin{aligned} var(X) &= E \left((X - E(X))^2 \right) \\ &= E(X^2) - E(X)^2 \end{aligned}$$

- We use the notation $\sigma^2 = \sigma_X^2 = var(X)$
 σ : sigma, Σ : capitalized sigma
- The standard deviation is the square root of the variance: $\sigma = \sigma_X = \sqrt{var(X)}$

Variance

- Recall: $\text{var}(X) = E \left((X - E(X))^2 \right)$
- Let X be a discrete random variable with mean μ_X :

$$\sigma_X^2 = \sum_{S_X} (x - \mu_X)^2 \Pr(X = x)$$

- Let X be a continuous random variable with mean μ_X

$$\sigma_X^2 = \int_{-\infty}^{\infty} (x - \mu_X)^2 f_X(x) dx$$

Variance: Example

- Setup: Flip two fair coins; let X be the number of heads
- Q: What is $\text{var}(X)$?
- Q: What is the standard deviation of X ?

Variance: Example

- Setup: Flip two fair coins; let X be the number of heads
- Q: What is $\text{var}(X)$?

We have that $\sigma_X^2 = \sum_{s_X} (x - \mu_X)^2 \Pr(X = x)$

We also know that $\mu_X = 1$, $P(X = 0) = P(X = 2) = 1/4$, and $P(X = 1) = 1/2$

Therefore, $\sigma_X^2 = (0 - 1)^2 \cdot (1/4) + (1 - 1)^2 \cdot (1/2) + (2 - 1)^2 \cdot (1/4) = 1/2$

- Q: What is the standard deviation of X ?

$$\sigma_X = \sqrt{\sigma_X^2} = \frac{1}{\sqrt{2}}$$

Functions of Random Variables

- Take random variable X and function $g(\cdot)$
- We can obtain a new random variable: $Y = g(X)$
- This is what is called a *transformation of variables*
- In general, to get the distribution of Y , we have that for any event $E \subseteq \mathcal{S}_Y$, we have $p_Y(E) = p_X(g^{-1}(E))$

如果 $Y = g(X)$ ，事件 Y 的发生概率和与其对应的事件 X 的发生概率相等

Linear Transformations: Mean and Variance

- Let g be a linear function of the form $g(X) = aX + b$
- Let X be a random variable with mean μ_X and variance σ_X^2
- Define a new random variable $Y = g(X) = aX + b$
- Finding the mean of Y :

$$\mu_Y = E(Y) = E(aX + b) = aE(X) + b = a\mu_X + b$$

- Finding the variance of Y :

$$\begin{aligned}\sigma_Y^2 &= \text{var}(Y) = E((aX + b - E(aX + b))^2) \\ &= E((aX + b - aE(X) - b)^2) = E((aX - aE(X))^2) \\ &= E(a^2(X - E(X))^2) = a^2E((X - E(X))^2) \\ &= a^2 \cdot \text{var}(X) = a^2 \cdot \sigma_X^2\end{aligned}$$

General Transformations: Mean

- If we have $Y = g(X)$ for general $g(X)$, then we have:

$$\mu_Y = E(Y) = E(g(X))$$

- We do **not** necessarily have that:

$$E(g(X)) = g(E(X))$$

- Example: Consider X = the outcome of rolling a fair six-sided die, and let $g(X) = X^2$

$$E(g(X)) = \frac{1^2 + 2^2 + 3^2 + 4^2 + 5^2 + 6^2}{6} = \frac{91}{6} \approx 15.17$$

$$g(E(X)) = \left(\frac{1 + 2 + 3 + 4 + 5 + 6}{6} \right)^2 = (3.5)^2 = 12.25$$

Independence

- Two random variables X_1 and X_2 are independent if the following holds, for any two events E_1 and E_2 :

$$P(X_1 \in E_1 \cap X_2 \in E_2) = P(X_1 \in E_1) \cdot P(X_2 \in E_2)$$

- Notation:

$X_1 \perp X_2$ means X_1 and X_2 are independent

- If a collection of random variables X_1, X_2, \dots, X_n are all independent and have the same distribution, we say that they are i.i.d. (independent and identically distributed)
 - Example: Roll two dice, or flip three fair coins

Covariance

- If two variables are not independent, we measure their dependency through their **covariance**
- Let X and Y be two random variables with means μ_X and μ_Y , respectively
- The covariance of X and Y is defined as follows:
 $\mu_X = E(X)$, $\mu_Y = E(Y)$, they can be considered as constants

$$\text{cov}(X, Y) = \sigma_{XY} = E((X - \mu_X)(Y - \mu_Y)) = E(XY) - E(X)E(Y)$$

- Correlation (essentially normalized covariance):

$$\text{corr}(X, Y) = \rho = \rho_{XY} = \frac{\text{cov}(X, Y)}{\sigma_X \sigma_Y}$$

Properties of Covariance

- Given random variables X and Y , the following hold:
 - If either X or Y is a constant, then $cov(X, Y) = 0$ and $corr(X, Y)$ is undefined
 - If $X \perp Y$, then $cov(X, Y) = corr(X, Y) = 0$
 - $cov(X, X) = var(X)$
 - $cov(X, Y) = cov(Y, X)$

Linear Combinations

- Suppose you have random variables X and Y with means μ_X, μ_Y and variances σ_X^2, σ_Y^2
- Let $Z = aX + bY$
- The mean of Z is

$$\mu_Z = E(Z) = E(aX + bY) = E(aX) + E(bY) = a\mu_X + b\mu_Y$$

- The variance of Z is

$$\sigma_Z^2 = a^2\sigma_X^2 + b^2\sigma_Y^2 + 2ab\sigma_{XY}$$

- The standard deviation of Z is

$$\sigma_Z = \sqrt{a^2\sigma_X^2 + b^2\sigma_Y^2 + 2ab\sigma_{XY}}$$

Theoretical Distributions

- Theoretical probability distributions describe what we expect to happen based on populations on a theoretical level
- We will consider the following theoretical distributions (D = discrete, C = continuous):
 - Bernoulli distribution (D)
 - Binomial distribution (D)
 - Poisson distribution (D)
 - Geometric distribution (D)
 - Uniform distribution (C)
 - Exponential distribution (C)
 - Normal distribution (C)

Bernoulli Distribution

- Let Y be a dichotomous random variable (takes one of two mutually exclusive values)
 - Classic example: Coin flip
- Successes ($= 1$) occur with probability p and failures ($= 0$) occur with probability $1 - p$, for constant $p \in [0,1]$
- Notation: $Y \sim \text{Bern}(p)$

Bernoulli Distribution: Example

- Let Y be a dichotomous random variable representing a coin flip
 - $Y = 1$: heads
 - $Y = 0$: tails
- If the coin is fair, then $p =$ and $1 - p =$
- If the coin has a 60% chance of landing heads, then $p =$ and $1 - p =$
- What is $E(Y)$ in terms of p ?
- What is $var(Y)$ in terms of p ?


Bernoulli Distribution: Example

- Let Y be a dichotomous random variable representing a coin flip
 - $Y = 1$: heads
 - $Y = 0$: tails
- If the coin is fair, then $p = 1/2$ and $1 - p = 1/2$
- If the coin has a 60% chance of landing heads, then $p = 0.6$ and $1 - p = 0.4$
- What is $E(Y)$ in terms of p ? $E(Y) = p \cdot 1 + (1 - p) \cdot 0 = p$
- What is $var(Y)$ in terms of p ? $var(Y) = E(Y^2) - E(Y)^2 = p - p^2 = p(1 - p)$

Binomial Distribution

- Suppose we flip n i.i.d. coins instead of just one coin
- Let $X = \sum_{i=1}^n X_i$ be the number of heads we observe
- X is binomially distributed
- **Binomial distribution:** If we have a sequence of n Bernoulli random variables, each with a probability of success p , then the total number of successes is a binomial random variable
 - Assumptions: fixed number of trials, independent, constant p
- Notation: $X \sim \text{Bin}(n, p)$

Binomial Distribution



Binomial Coefficients

- Let $X = \sum_{i=1}^n X_i$ be the number of heads we observe when we flip n i.i.d. coins
- Each coin has probability of heads p , and flips are independent
- Q: What is the probability of getting exactly x out of n successes?
 - Choose which x trials succeed:
 - Probability that these x trials succeed:
 - Probability that the other $n - x$ trials fail:
- In general,

Binomial Coefficients

- Let $X = \sum_{i=1}^n X_i$ be the number of heads we observe when we flip n i.i.d. coins
- Each coin has probability of heads p , and flips are independent
- Q: What is the probability of getting exactly x out of n successes?
 - Choose which x trials succeed: $\binom{n}{x}$
 - Probability that these x trials succeed: p^x
 - Probability that the other $n - x$ trials fail: $(1 - p)^{n-x}$
- In general, $\Pr(X = x) = \binom{n}{x} p^x (1 - p)^{n-x}$

Binomial Probabilities in R

- Calculate probabilities in R:
 - Calculate $\Pr(X = x)$ using `dbinom(x, n, p)`
 - Calculate $\Pr(X \leq x)$ using `pbinom(x, n, p)`
 - Calculate $\Pr(X \geq x)$ using `1-pbinom(x-1, n, p)`

Binomial Distribution: Summary Measures

- Note that a binomial distribution with parameters n and p is the sum of n independent Bernoulli distributions with parameter p

$$E(X) = \mu_X = np$$

$$\text{var}(X) = \sigma_X^2 = np(1 - p)$$

$$\text{stdev}(X) = \sigma_X = \sqrt{np(1 - p)}$$

- Q: How does $\text{var}(X)$ change with $p \in [0,1]$?

Binomial Distribution: Summary Measures

- Note that a binomial distribution with parameters n and p is the sum of n independent Bernoulli distributions with parameter p

$$E(X) = \mu_X = np$$

$$\text{var}(X) = \sigma_X^2 = np(1 - p)$$

$$\text{stdev}(X) = \sigma_X = \sqrt{np(1 - p)}$$

- Q: How does $\text{var}(X)$ change with $p \in [0,1]$?
Highest when $p = 0.5$, lowest (= 0) when $p = 0$ or 1

Binomial Distribution: Summary

- Main take-away points from the binomial distribution:
 - Fixed number of independent Bernoulli trials, n
 - Constant probability of success, p (Bernoulli parameter)
 - Interested in the total number of successes in n trials (not order)
 - Mean: $\mu_X = np$
 - Variance: $\sigma_X^2 = np(1 - p)$
- Examples:
 - Number of heads in 15 flips of a fair coin

Poisson Distribution

- **Poisson distribution:** Probability of observing a certain number of discrete events within a known interval
 - Models discrete events that occur infrequently in time or space
- Example:
 - The number of babies born at Strong Memorial Hospital between 10 am and 2 pm today
 - The number of students who enter River Campus today

Poisson Distribution

- Let $X \in [0, \infty)$ be the number of occurrences of some event over a given interval
- Let $\lambda > 0$ be the average number of occurrences of the event over the specified interval
- In this case, we say that $X \sim \text{Pois}(\lambda)$
- The probability function is given by $P(X = x) = \frac{e^{-\lambda} \lambda^x}{x!}$
- If $X \sim \text{Pois}(\lambda)$, then $\mu_X = \sigma_X^2 = \lambda$
 - For a Poisson distribution, both the mean and the variance are equal to λ

Poisson Distribution

- Poisson distribution assumptions:
 - The probability of an event occurring is proportional to the length of the interval
 - Within an interval, an infinite number of events is theoretically possible
 - Events occur independently
 - The number of events that occur must be non-negative

Poisson Distribution: Example

- Setup: We want to examine the probability of certain numbers of people developing a rare disease in the next year. On average, 1.95 people develop the disease per year
- Q: What is the probability of no one developing the disease in the next year?
- Q: What is the probability of one person developing the disease in the next year?

Poisson Distribution: Example


- Setup: We want to examine the probability of certain numbers of people developing a rare disease in the next year. On average, 1.95 people develop the disease per year
- Q: What is the probability of no one developing the disease in the next year?

$$P(X = 0) = \frac{e^{-1.95} 1.95^0}{0!} = e^{-1.95} \approx 0.142$$

- Q: What is the probability of one person developing the disease in the next year?

$$P(X = 1) = \frac{e^{-1.95} 1.95^1}{1!} = 1.95 \cdot e^{-1.95} \approx 0.277$$

Poisson Distribution: Visualized



Poisson Probabilities in R

- Calculate probabilities in R:
 - Calculate $\Pr(X = x)$ using `dpois(x, lambda)`
 - Calculate $\Pr(X \leq x)$ using `ppois(x, lambda)`
 - Calculate $\Pr(X \geq x)$ using `1-ppois(x-1, lambda)`

Poisson Distribution: Summary

- Main take-away points from the Poisson distribution:
 - Fixed interval, independent events, interested in number of events in interval
 - Unlimited number of events is theoretically possible
 - Mean: $\mu_X = \lambda$
 - Variance: $\sigma_X^2 = \lambda$
- Examples:
 - Number of calculators the book store sells this year
 - Number of babies born today

Geometric Distribution

- Suppose Y_1, Y_2, \dots is an *infinite* sequence of independent Bernoulli random variables with parameter p
- Let X be the first index i for which $Y_i = 1$ (location of first success)
- The probability mass function (PMF) is given by

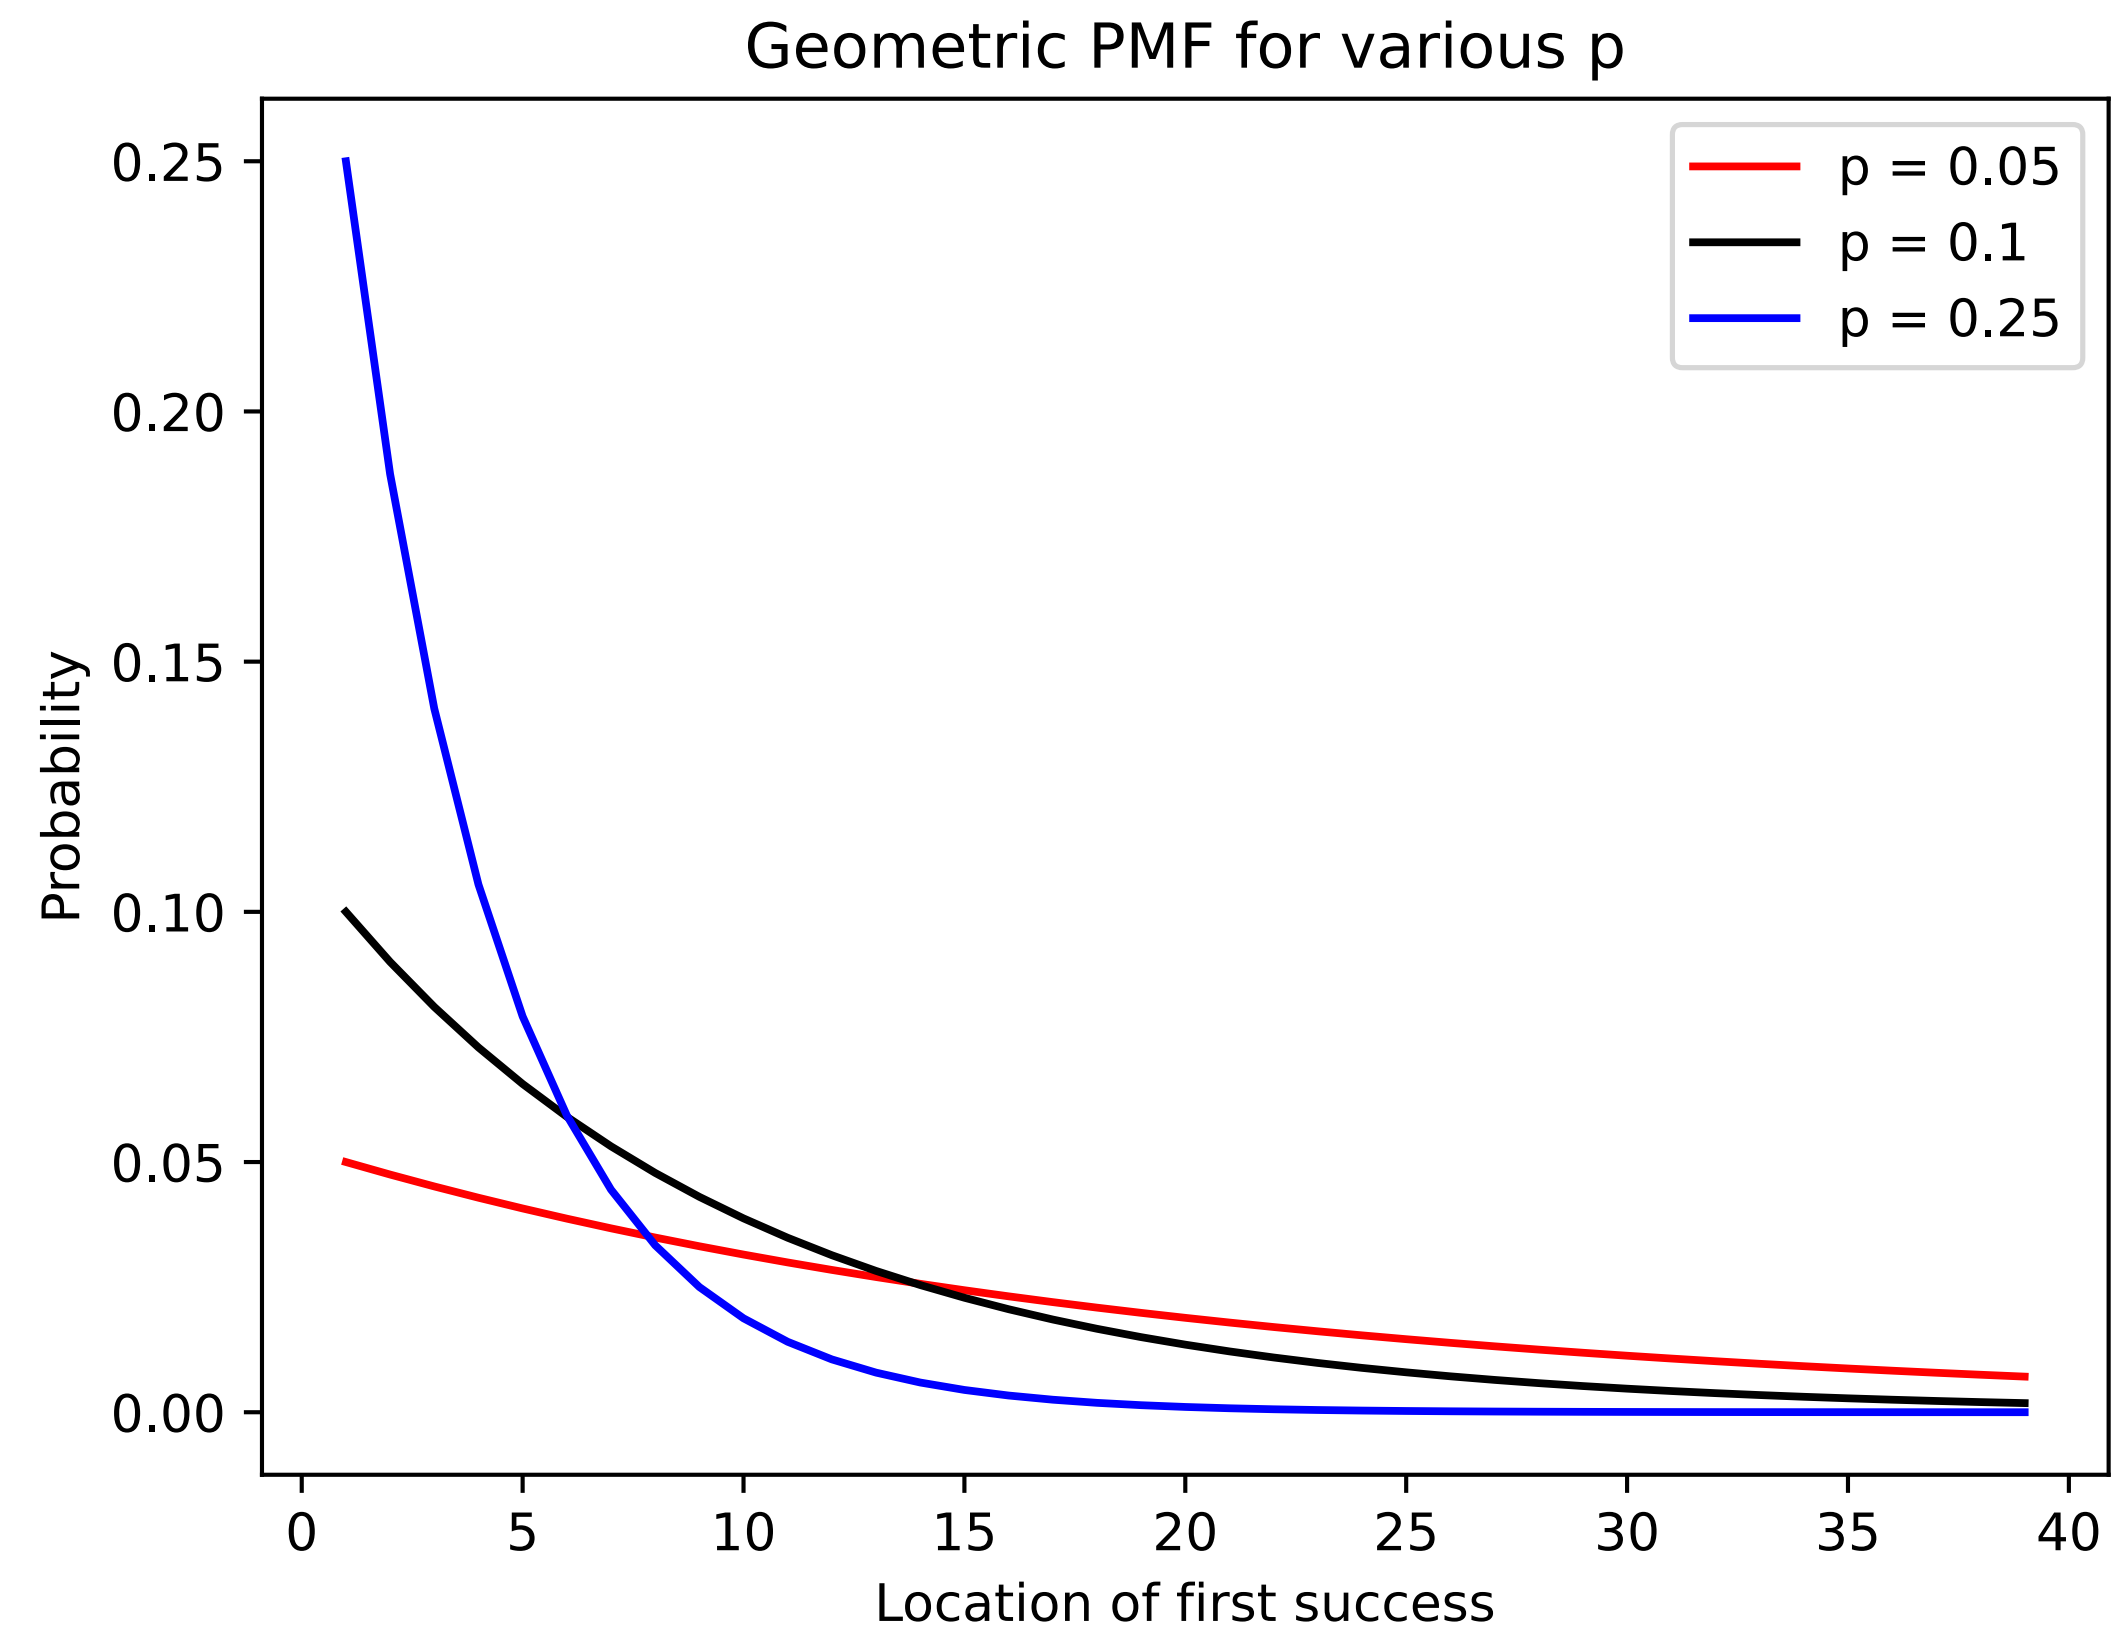
$$P(X = x) = p(1 - p)^{x-1}$$

- Notation: $X \sim \text{Geom}(p)$
- Mean and variance:

$$E(X) = \mu_X = \frac{1}{p}$$

$$\text{var}(X) = \sigma_X^2 = \frac{1 - p}{p^2}$$

- CDF: $P(X \leq x) = 1 - (1 - p)^x$



Continuous Distributions

- Continuous random variables: Intuitively, discrete random variables that are infinitesimally close together
- Instead of having discrete bars for the density at each discrete value, we now have a continuous density curve
- The area under the curve equals 1 (law of total probability)
- Since the random variable X can take on an infinite number of values, the probability associated with any single outcome equals 0
- The probability that $X \in (x_1, x_2)$ is equal to the area under the curve that lies between these two values


Uniform Distribution

- Let X be a continuous random variable which can take on any value between a and b with equal probability
- Any value outside the range $[a, b]$ cannot occur
- The uniform probability density is

$$f_X(x) = \begin{cases} \frac{1}{b-a} & x \in [a, b] \\ 0 & \text{otherwise} \end{cases}$$

- Notation: $X \sim \text{Unif}(a, b)$

- $\mu_X = \frac{a+b}{2}$ and $\sigma_X^2 = \frac{(b-a)^2}{12}$



Exponential Distribution


- A continuous random variable X is exponentially distributed if it follows the following density:

$$f_X(x) = \lambda e^{-\lambda x}, \lambda > 0$$

- Notation: $X \sim \text{Exp}(\lambda)$
- Generalizes the geometric distribution

- $\mu_X = \frac{1}{\lambda}$ and $\sigma_X^2 = \frac{1}{\lambda^2}$

- CDF: $F_X(x) = 1 - e^{-\lambda x}$



Normal Distribution


- The most common continuous distribution is the normal distribution (also called a Gaussian distribution or bell-shaped curve)
 - Shape of the binomial distribution when p is constant but $n \rightarrow \infty$
 - Shape of the Poisson distribution when $\lambda \rightarrow \infty$
- Given μ, σ , the density function is

$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$

- Notation: $X \sim N(\mu, \sigma^2)$ – but in R, use standard deviation instead of variance
- Mean = median = mode = μ , variance = σ^2 , standard deviation = σ

Normal Distribution: Visualization

- μ (center) and σ^2 (spread) fully define the normal distribution
- Always symmetric
- When $\mu = 0$ and $\sigma^2 = 1$, we have the *standard normal distribution*



Normal Distribution: z-scores

- Recall from Chapter 2 that a z-score tells us how many standard deviations an observation is from its mean:

$$z = \frac{x - \mu}{\sigma}$$

- Z has the nice property that it will always be $N(0,1)$
- Given $X \sim N(\mu, \sigma)$, we can calculate a z-score, which will be $Z \sim N(0,1)$
- Standardizes the procedures for all normal distribution problems

Normal Distribution

- Recall that for continuous distributions, we are interested in determining the probability of being in an interval:

$$\Pr(X \leq a), \Pr(X \geq b), \text{ or } \Pr(a \leq X \leq b)$$

- We can look at the plot of the normal distribution and determine the probability (= area under the curve between endpoints)
- In general, we will use R to calculate area under the curve (i.e., probabilities)
- By default, R works in terms of z-scores:

$$\Pr(Z \leq z) : \text{pnorm}(z)$$

$$\Pr(Z \geq z) : 1 - \text{pnorm}(z)$$

$$\Pr(z_1 \leq Z \leq z_2) : \text{pnorm}(z_2) - \text{pnorm}(z_1)$$

Normal Distribution

- The general process for calculating probabilities based on a normal distribution is as follows:
 - Calculate appropriate z-scores: $z = \frac{X - \mu}{\sigma}$
 - Use R to calculate the probability based on this z-score (`pnorm(z)`)

Normal Distribution: Example

- Suppose that test scores are normally distributed with mean 78 and standard deviation 9
- Q: What is the probability that a person scored below 60?
- Q: What is the probability that a person scored between 80 and 90?

Normal Distribution: Example

- Suppose that test scores are normally distributed with mean 78 and standard deviation 9

- Q: What is the probability that a person scored below 60?

$$\text{z-score: } \frac{60 - 78}{9} = -2$$

Use R: `pnorm(-2) = 0.0228`

- Q: What is the probability that a person scored between 80 and 90?

$$\text{z-scores: } \frac{80 - 78}{9} = 2/9, \frac{90 - 78}{9} = 4/3$$

Use R: `pnorm(4/3) - pnorm(2/9) = 0.3209`

Normal Probabilities in R (Shortcut)

- We can let R do the entire process of calculating a z-score and probability for us
- Let $X \sim N(\mu, \sigma)$

$$\Pr(X \leq x) : \text{pnorm}(x, \text{mean}, \text{sd})$$

$$\Pr(X \geq x) : 1 - \text{pnorm}(x, \text{mean}, \text{sd})$$

$$\Pr(x_1 \leq X \leq x_2) : \text{pnorm}(x_2, \text{mean}, \text{sd}) - \text{pnorm}(x_1, \text{mean}, \text{sd})$$

Normal Distribution

- Suppose that BMI is normally distributed with mean 26.6 and standard deviation 3.2
- What is the probability that a person has a BMI in the range (18.5, 24.9)?
- Find $\Pr(18.5 \leq X \leq 24.9)$:

$$z_1 = \frac{18.5 - 26.6}{3.2} = -2.53$$

$$z_2 = \frac{24.9 - 26.6}{3.2} = -0.53$$

$$\Pr(-2.53 \leq Z \leq -0.53) = 0.2924$$

- $\text{pnorm}(-0.53) - \text{pnorm}(-2.53) = 0.2924$
- $\text{pnorm}(24.9, 26.6, 3.2) - \text{pnorm}(18.5, 26.6, 3.2) = 0.2919$

Revisiting the Empirical Rule

- How well does the empirical rule approximate the normal distribution?
 - $\Pr(-1 \leq Z \leq 1) = 0.683$
 - $\Pr(-2 \leq Z \leq 2) = 0.954$
 - $\Pr(-3 \leq Z \leq 3) = 0.997$
- Empirical rule (68%, 95%, 99.7%) appears to be quite good

Normal Distribution: Percentiles

- Given data x_1, \dots, x_n , what value of x corresponds to a probability of the p^{th} percentile?
- Strategy:
 - Find z value such that $\Pr(Z \leq z) = p$ (lower tail probability is p)
 - Solve for x by inverting z-score: $x = z \cdot \sigma + \mu$
- Directly in R: `qnorm(p, mean, sd)` ; `qnorm(p)` for z value

Normal Distribution: Example

- Setup: Let X be a random variable that represents weights of patients in American hospital EDs; X is normally distributed with $\mu = 160$ and $\sigma = 15$
- Q1: Find the probability that a randomly selected patient in the ED weighs between 140 pounds and 210 pounds
- Q2: Find the value that cuts off the upper 10% of the curve in American ED patient weights

Normal Distribution: Example

- Setup: Let X be a random variable that represents weights of patients in American hospital EDs; X is normally distributed with $\mu = 160$ and $\sigma = 15$
- Q1: Find the probability that a randomly selected patient in the ED weighs between 140 pounds and 210 pounds

Find z-scores: $z = \frac{x - \mu}{\sigma}$, so $z_1 = \frac{140 - 160}{15} = -4/3$ and $z_2 = \frac{190 - 160}{15} = 2$

$$\text{pnorm}(2) - \text{pnorm}(-4/3) = 0.886$$

- Q2: Find the value that cuts off the upper 10% of the curve in American ED patient weights

Find z-score: $z_{0.9} = \text{qnorm}(0.9) = 1.282 = \frac{x - 160}{15}$

$$x = 160 + 1.282 \cdot 15 = 179.2$$

Sampling Distributions

- Suppose we want to estimate the mean value of some continuous random variable of interest
- We can take a sample from the population and use the sample mean as an estimate of the population mean: \bar{x} is an estimate for μ
- For a normally distributed population, \bar{x} is the *maximum likelihood estimator* for μ
 - Value of the parameter that is most likely to have produced the observed sample data
- Different samples will have different means

Sampling Distributions

- What if you continued sampling m times?
 - You take one random sample and get mean \bar{x}_1 , take another random sample and get mean \bar{x}_2 , and repeat until you have $\bar{x}_1, \bar{x}_2, \dots, \bar{x}_m$
 - Take m random samples for a total of m sample means
- These m means form a distribution with mean μ and variance $\frac{\sigma^2}{n}$ where n is the sample size
- Key idea: \bar{X} has its own distribution
- Standard deviation of \bar{X} is $\frac{\sigma}{\sqrt{n}}$; this is known as the **standard error**

Sampling Normal Distributions

- If the population we are sampling from is normal, then the distribution of \bar{X} will also be normal
- If $X \sim N(\mu, \sigma)$, then $\bar{X} \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$

Central Limit Theorem (CLT)

- If the population we are sampling from is not normal, then we can use the *Central Limit Theorem (CLT)* to get the distribution of \bar{X}
- **Central Limit Theorem:** If the population we are sampling from is not normal, then the shape of the distribution of \bar{X} will be normal as long as n is sufficiently large (typically $n \geq 30$ suffices)

Central Limit Theorem (CLT)

- In particular, given that the distribution of an underlying population has mean μ and standard deviation σ , the distribution of the sample means computed for samples of size n has three important properties:
 - The mean of the sampling distribution equals the population mean μ
 - The standard deviation of the distribution of sample means is equal to $\frac{\sigma}{\sqrt{n}}$, which is the standard error of the mean
 - Given that n is sufficiently large, the shape of the sampling distribution is approximately normal
- In notation: $\bar{X} \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$

Central Limit Theorem (CLT)

- The further the underlying population is from normal, the larger the sample size you need to ensure normality of the sampling distribution
- However, if the underlying population is normal, you do not need the central limit theorem to ensure normality of the sampling distribution – normality will hold regardless of the sample size if the underlying population is normal
- Since $\bar{X} \sim N\left(\mu, \sigma/\sqrt{n}\right)$, we can standardize \bar{X} to a standard normal distribution as follows:

$$Z = \frac{\bar{X} - \mu}{\frac{\sigma}{\sqrt{n}}}$$

Sampling Distributions Example

- Setup: Suppose house prices have a distribution with a mean of $\mu = \$450,000$ and standard deviation $\sigma = \$100,000$. You draw a random sample of $n = 64$ houses and determine their prices
- Q1: What is the mean of the distribution of sample means?
- Q2: What is the standard error of the sample mean?
- Q3: What distribution does the sample mean follow?
- Q4: What is the probability that the sample mean of $n = 64$ house prices is greater than \$500,000?

Sampling Distributions Example

- Setup: Suppose house prices have a distribution with a mean of $\mu = \$450,000$ and standard deviation $\sigma = \$100,000$. You draw a random sample of $n = 64$ houses and determine their prices
- Q1: What is the mean of the distribution of sample means?
 $\mu = \$450,000$
- Q2: What is the standard error of the sample mean?
 $\sigma/\sqrt{n} = \$100,000/8 = \$12,500$
- Q3: What distribution does the sample mean follow?
By CLT, $\bar{X} \sim N(450000, 12500)$
- Q4: What is the probability that the sample mean of $n = 64$ house prices is greater than \$500,000?
 $\Pr(\bar{X} > 500000) = \Pr\left(Z > \frac{500000 - 450000}{12500}\right) = \Pr(Z > 4) = 0.000032$

Sampling Distribution of a Proportion

- Suppose we are interested in the proportion of the time that an event occurs
- If we take a sample of size n and observe x successes, then we could estimate the population proportion p by $\hat{p} = x/n$
- We can do this sampling process m times for a total of m different values of \hat{p}_i for $i \in \{1, 2, \dots, m\}$
- These m proportions form a distribution with mean p
- Standard deviation of \hat{p} , known as standard error, is $\sqrt{\frac{p(1-p)}{n}}$

Sampling Distribution of a Proportion

- The shape of the distribution of \hat{p} will be approximately normal as long as two conditions are met:
 - $np \geq 5$
 - $n(1 - p) \geq 5$
- If both of these conditions are met, then $\hat{p} \sim N\left(p, \sqrt{\frac{p(1 - p)}{n}}\right)$

Sampling Distribution of a Proportion: Example

- Setup: Suppose 20% of Americans favor Advil as a pain reducer. A polling organization takes a sample of 100 Americans and asks if they prefer Advil or some other pain relief medicine.

- Q1: What is the mean of this sample proportion?

$$\mu = 0.20$$

- Q2: What is the standard error of this sample proportion?

$$\sqrt{\frac{0.2(1 - 0.2)}{100}} = 0.04$$

- Q3: What distribution does the sample proportion follow?

$$np = 20 > 5, \text{ and } n(1 - p) = 80 > 5, \text{ so by CLT, } \hat{p} \sim N(0.2, 0.04)$$

- Q4: What is the probability that the sample proportion is less than 18%?

$$\Pr(\hat{p} < 0.18) = \Pr(Z < (0.18 - 0.2)/0.04) = \Pr(Z < -0.5) \approx 0.31$$

- Q5: What is the 20th percentile of the distribution of the sample proportion?

$$z_{0.20} = \frac{x - \mu}{\sigma} \rightarrow x = 0.2 + (-0.84) \cdot 0.04 \approx 0.167$$