

Chapter 10: Inference on Proportions

DSCC 462

Computational Introduction to Statistics

Anson Kahng

Fall 2022

Inference on Proportions

- So far, we have considered inference for when we have continuous data
- We can also extend inferential methods to cover count data
- In particular, we are often interested in the proportion of times a dichotomous (i.e., yes/no) event occurs

Sampling Distribution of a Proportion

- Recall that the sample mean is distributed like $\hat{p} \sim N\left(p, \sqrt{\frac{p(1-p)}{n}}\right)$,
given that $np \geq 5$ and $n(1-p) \geq 5$
- Thus, $Z = \frac{\hat{p} - p}{\sqrt{\frac{p(1-p)}{n}}}$ has a standard normal distribution

Confidence Intervals for Proportions

- Confidence intervals for population proportions follows the same procedure as what we used for population means
- Draw a sample of size n and compute $\hat{p} = \frac{x}{n}$
- \hat{p} is a point estimate of population proportion p
- We know from above that $Z = \frac{\hat{p} - p}{\sqrt{\frac{p(1-p)}{n}}}$ is a standard normal random variable, given that n is sufficiently large (i.e., $np \geq 5$ and $n(1-p) \geq 5$)

Confidence Intervals for Proportions

- For a standard normal distribution, 95% of possible outcomes lie between $qnorm(0.025) = -1.96$ and $qnorm(0.975) = 1.96$

- Thus,
$$\Pr \left(-1.96 \leq \frac{\hat{p} - p}{\sqrt{\frac{p(1-p)}{n}}} \leq 1.96 \right) = 0.95$$

- This can be rearranged to give

$$\Pr \left(\hat{p} - 1.96 \sqrt{\frac{p(1-p)}{n}} \leq p \leq \hat{p} + 1.96 \sqrt{\frac{p(1-p)}{n}} \right) = 0.95$$

- Note that this confidence interval depends on the (unknown) value of p !

Confidence Intervals for Proportions

- So how do we estimate p ? Use \hat{p} , our sample estimate (Wald)

- Therefore, our confidence interval calculation becomes

$$\Pr \left(\hat{p} - 1.96\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \leq p \leq \hat{p} + 1.96\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right) = 0.95$$

- In other words, we are 95% confident that the interval

$$\left(\hat{p} - 1.96\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, \hat{p} + 1.96\sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right) \text{ contains the true population}$$

proportion p

Confidence Intervals for Proportions

- In general, an approximate two-sided $(1 - \alpha) \cdot 100\%$ confidence interval for p is given by $\left(\hat{p} - z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, \hat{p} + z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right)$
- A one-sided lower $(1 - \alpha) \cdot 100\%$ confidence interval for p is given by $\left(\hat{p} - z_{\alpha} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, 1 \right)$
- A one-sided upper $(1 - \alpha) \cdot 100\%$ confidence interval for p is given by $\left(0, \hat{p} + z_{\alpha} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right)$

Wald vs. Wilson Intervals

- Recall that we estimate p using \hat{p} , our sample estimate (Wald)
- In general, this provides poor coverage when \hat{p} is close to extremes (0 or 1)
 - Less than $(1 - \alpha) \cdot 100\%$ confidence interval
- (One) alternative method: Wilson (what `prop.test()` uses in R)
 - Solve for p in terms of Z, \hat{p} from the approximation $Z = \frac{\hat{p} - p}{\sqrt{\frac{p(1-p)}{n}}}$
 - Get an estimate of p that is a weighted average of \hat{p} and $\frac{1}{2}$, where the weight on \hat{p} increases with n
 - Better coverage!

Confidence Intervals for Proportions: Example

- Setup: We are interested in determining what proportion of a population is right-handed. Suppose we have a sample of $n = 62$ subjects and 53 of these subjects are right-handed. Find a 95% confidence interval for the population proportion of right-handed people in the population
- Find \hat{p} (our estimate of p):
- Check normality assumptions:
- Apply a two-sided 95% confidence interval:

Confidence Intervals for Proportions: Example

- Setup: We are interested in determining what proportion of a population is right-handed. Suppose we have a sample of $n = 62$ subjects and 53 of these subjects are right-handed. Find a 95% confidence interval for the population proportion of right-handed people in the population
- Find \hat{p} (our estimate of p): $\hat{p} = 53/62 = 0.855$
- Check normality assumptions: $n\hat{p} = 62(0.855) = 53 \geq 5$ and $n(1 - \hat{p}) = 62(1 - 0.855) = 8.99 \geq 5$, so normal approximation is appropriate
- Apply a two-sided 95% confidence interval: $\left(\hat{p} - z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}, \hat{p} + z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \right)$, plugging in yields (0.767, 0.943)
- We are 95% confident that the interval (0.767, 0.943) contains the true population proportion of people who are right-handed

Normal Approximations of Binomial Distributions

- Note that the true distribution for proportions is a binomial distribution (number of “successes” out of a certain number of trials)
- However, confidence intervals are based on the normal distribution
- We are using the normal distribution as an approximation for a binomial distribution
- Normal approximation (Wilson) confidence intervals can be calculated in R using `prop.test(x, n)`
- Exact binomial (Clopper-Pearson) confidence intervals can be calculated in R using `binom.test(x, n)`

Sample Size Estimation

- For confidence intervals on proportions, we have that the margin of error is

$$m = z_{\alpha/2} \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}} \text{ (half the length of the confidence interval)}$$

- Just like before, if we want a certain margin of error at the same confidence level, we can determine the number of subjects (n) needed to get the desired results

- Thus,
$$n = \left\lceil \frac{z_{\alpha/2}^2 p(1 - p)}{m^2} \right\rceil$$

- If we can estimate p based on previous studies or information, use that
- Otherwise, use $p = 0.5$ to get the most conservative estimate of the standard error (overestimate of the number of subjects needed)

Sample Size Estimation: Example

- Setup: We want to determine what proportion of college students have an iPhone within a margin of error of 8 percentage points with 95% confidence
- Q1: How large of a sample should you take?
- Q2: A national study determined that 38% of all Americans own iPhones. Now, how large of a sample should you take?

Sample Size Estimation: Example

- Setup: We want to determine what proportion of college students have an iPhone within a margin of error of 8 percentage points with 95% confidence
- Q1: How large of a sample should you take?

$$m = 0.08, \text{ and we assume } p = 0.5, \text{ so } n = \left\lceil \frac{1.96^2 \cdot 0.5 \cdot (1 - 0.5)}{0.08^2} \right\rceil = 151 \text{ people}$$

- Q2: A national study determined that 38% of all Americans own iPhones. Now, how large of a sample should you take?

$$m = 0.08 \text{ and we have an estimate of } p = 0.38, \text{ so } n = \left\lceil \frac{1.96^2 \cdot 0.38 \cdot (1 - 0.38)}{0.08^2} \right\rceil = 142 \text{ people}$$

Hypothesis Testing for Proportions

- Just as we used hypothesis tests to see if a population mean was equal to some hypothesized value, we can also test whether a population proportion is equal to some value
- Consider a two-tailed test at the $\alpha = 0.05$ significance level
- $H_0 : p = p_0$ vs. $H_1 : p \neq p_0$
- Draw a random sample of size n observations from the underlying population (each observation is a dichotomous yes/no)

- Calculate a z-statistic:
$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}}$$

Hypothesis Testing for Proportions

- For sufficiently large n and when H_0 is true, we can compare z to a standard normal distribution to calculate the probability of obtaining a proportion as extreme or more extreme than \hat{p}
- Calculate p-value in R by `p=2*pnorm(-abs(z))`
- If $p \leq 0.05$, we reject the null hypothesis and conclude that $p \neq p_0$
- If $p > 0.05$, we fail to reject the null hypothesis and conclude that there is not significant evidence to say that $p \neq p_0$

Hypothesis Testing for Proportions: Example

- Consider the dominant hand example ($n = 62$, $\hat{p} = 53/62 = 0.855$)
- Test at the $\alpha = 0.05$ level whether the true population proportion of right-handed people is equal to 0.9
- $H_0 : p = 0.9$ vs. $H_1 : p \neq 0.9$
- Check normality assumptions based on p :
- Calculate z-score:
- Calculate p-value:

Hypothesis Testing for Proportions: Example

- Consider the dominant hand example ($n = 62$, $\hat{p} = 53/62 = 0.855$)
- Test at the $\alpha = 0.05$ level whether the true population proportion of right-handed people is equal to 0.9
- $H_0 : p = 0.9$ vs. $H_1 : p \neq 0.9$
- Check normality assumptions based on p : $np = 55.8$, $n(1 - p) = 6.2$
- Calculate z-score: $z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1 - p_0)}{n}}} = \frac{0.855 - 0.9}{\sqrt{\frac{0.9(1 - 0.9)}{62}}} = -1.18$
- Calculate p-value: $2 * \text{pnorm}(-1.18) = 0.238$
- Since $0.238 > 0.05$, we fail to reject the null hypothesis (insufficient evidence to conclude that the proportion of people who are right handed is different from 0.9)

Confidence Intervals vs. Hypothesis Tests for Proportions

- When looking at sample means, confidence intervals and hypothesis tests are essentially equivalent
- This is no longer the case for proportions!
 - Intuition: For sample means, there are two parameters of interest (μ , σ), whereas for proportions, p determines both the mean and variance
- For proportion hypothesis tests, we calculate the standard error based on p_0 as $\sqrt{\frac{p_0(1-p_0)}{n}}$ (i.e., our frame of reference is centered at the null hypothesis)
- For proportion (Wald) confidence intervals, we calculate the standard error based on \hat{p} as $\sqrt{\frac{\hat{p}(1-\hat{p})}{n}}$ (i.e., our frame of reference is centered at our observed sample proportion)
- As with confidence intervals, we can perform an exact test based on the binomial distribution instead of using the normal approximation: `binom.test(x, n, p=p0)`

One-Sided Hypothesis Tests

- With two-sided hypothesis tests, we were only concerned with whether or not there was a difference from the postulated population proportion
 - $H_1 : p \neq p_0$
- However, we are sometimes interested in deviations only in one direction
 - $H_1 : p > p_0$
 - $H_1 : p < p_0$
- For two-sided tests, we are concerned with the area in both tails of the distribution
- For one-sided tests, we are concerned with the area in only one tail of the distribution
- Analyses follow directly as they did for one-sided tests for sample means

Comparison of Two Proportions

- We can extend hypothesis tests to situations where we compare proportions for two groups
- Interested in testing whether the proportions from two independent populations are the same
 - $H_0 : p_1 = p_2$ or $H_0 : p_1 - p_2 = 0$
- Our alternative hypothesis is that there is a difference between these groups
 - $H_1 : p_1 \neq p_2$ or $H_1 : p_1 - p_2 \neq 0$

Comparison of Two Proportions

- We draw a sample of size n_1 from the first population and a sample of size n_2 from the second population
- There are x_1 successes in the first sample and x_2 successes in the second sample
- Sample proportion for each group:
 - $\hat{p}_1 = \frac{x_1}{n_1}$
 - $\hat{p}_2 = \frac{x_2}{n_2}$

Comparison of Two Proportions

- Under the null hypothesis, $p_1 = p_2 = p$
- Thus, the data from both samples can be combined to estimate this common parameter
 - $\hat{p} = \frac{x_1 + x_2}{n_1 + n_2} = \frac{n_1\hat{p}_1 + n_2\hat{p}_2}{n_1 + n_2}$
 - \hat{p} is the weighted average of the two sample proportions (or total successes over total trials)
- The estimator of the standard error of $\hat{p}_1 - \hat{p}_2$ can now be based on this common \hat{p}
 - Standard error: $\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}$
 - Similar to the “pooled” estimate for sample means

Comparison of Two Proportions

- Putting these pieces together, we get our z-statistic:

- $$z = \frac{(\hat{p}_1 - \hat{p}_2) - (p_1 - p_2)}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

- If n_1 and n_2 are sufficiently large, this z statistic is approximately a standard normal (mean 0, standard deviation 1)
- Typically, we want $n_1\hat{p}_1$, $n_1(1 - \hat{p}_1)$, $n_2\hat{p}_2$, and $n_2(1 - \hat{p}_2)$ to all be greater than 5 (this is a conservative standard)

Comparison of Two Proportions

- With these conditions satisfied, we compare the value of the z statistic with the critical value to find a p-value, p
- If $p \leq \alpha$, we reject the null hypothesis
- If $p > \alpha$, we fail to reject the null hypothesis

Comparison of Two Proportions

- Consider our dominant hand example. Suppose we are interested in knowing whether the right-handedness rate is different for Group A and Group B
- At the $\alpha = 0.01$ significance level, we will test the following hypotheses:
 - $H_0 : p_A = p_B$ vs. $H_1 : p_A \neq p_B$
- We take samples of $n_A = 54$ and $n_B = 62$
- We observe $x_A = 48$ and $x_B = 60$ subjects being right handed

Comparison of Two Proportions

- Calculate proportions:
- Is this difference too large to be attributed to chance?
- Under H_0 , $p_A = p_B$, so we can estimate their common value p

Comparison of Two Proportions

- Calculate proportions:

$$\hat{p}_A = \frac{x_A}{n_A} = \frac{48}{54} = 0.889$$

$$\hat{p}_B = \frac{x_B}{n_B} = \frac{60}{62} = 0.968$$

- Is this difference too large to be attributed to chance?
- Under H_0 , $p_A = p_B$, so we can estimate their common value p

$$\hat{p} = \frac{x_A + x_B}{n_A + n_B} = \frac{48 + 60}{54 + 62} = \frac{108}{116} = 0.931$$

Comparison of Two Proportions

- Checking normality assumptions:
 - $n_A p_A = 48 > 5$
 - $n_A(1 - p_A) = 6 > 5$
 - $n_B p_B = 60 > 5$
 - $n_B(1 - p_B) = 2 < 5 \implies$ proceed with caution

Comparison of Two Proportions

- Calculate z-statistic:
- Conclusion:

Comparison of Two Proportions

- Calculate z-statistic:

$$\begin{aligned} z &= \frac{\hat{p}_A - \hat{p}_B}{\sqrt{\hat{p}(1 - \hat{p})\left(\frac{1}{n_A} + \frac{1}{n_B}\right)}} \\ &= \frac{0.889 - 0.968}{\sqrt{0.931(1 - 0.931)\left(\frac{1}{54} + \frac{1}{62}\right)}} = -1.675 \\ \Rightarrow 2 \cdot \Pr(Z < -1.675) &= 0.09. \end{aligned}$$

- Since the p-value 0.09 is greater than $\alpha = 0.01$, we fail to reject the null hypothesis. There is not sufficient evidence to conclude that right-handedness differs between Group A and Group B

Comparison of Two Proportions

- We can also calculate a confidence interval for the difference of two proportions
- As in the one-sample case, the standard error is not the same for the confidence interval and hypothesis test
- For a two-sided confidence interval, we are $(1 - \alpha) \cdot 100\%$ confident that the interval $\left(\hat{p}_1 - \hat{p}_2 \pm z_{\alpha/2} \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}} \right)$ contains the true population difference, $p_1 - p_2$

Comparison of Two Proportions

- Continuing with our dominant hand by group example, we can construct a two-sided 95% confidence interval for $p_A - p_B$ as follows:

$$(0.889 - 0.968) \pm 1.96 \sqrt{\frac{0.889 \cdot 0.111}{54} + \frac{0.968 \cdot 0.032}{62}}$$
$$= (-0.173, 0.016)$$

- We are 95% confident that the interval $(-0.173, 0.016)$ contains the true difference in the proportion of members of Group A and Group B who are right-handed

Comparison of Two Proportions

- One tail, lower bound:

$$\left(\hat{p}_1 - \hat{p}_2 - z_\alpha \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}}, 1 \right)$$

- One tail, upper bound:

$$\left(-1, \hat{p}_1 - \hat{p}_2 + z_\alpha \sqrt{\frac{\hat{p}_1(1 - \hat{p}_1)}{n_1} + \frac{\hat{p}_2(1 - \hat{p}_2)}{n_2}} \right)$$

DSCC 462 Midpoint Survey

- <https://forms.gle/Zt3Qzrb7S7UXFXY28>
- If you fill it out: +2.5% on midterm
- If at least 90% of the class fills it out: +2.5% on each person's midterm
 - Tell your friends!